



Yann Maillot

APPLIED AI & FULL-STACK ENGINEER · RAG — LLM — AGENTS — ON-DEVICE ML

CONTACT

☎ +33 6 26 73 55 97
✉ yann.maillot.29@gmail.com
✦ Remote · anywhere in France
in /in/yann-maillot
gh @FelixDaCraft
➤ yann.aynn.fr

AI SKILLS

RAG: hybrid pipeline BGE-M3 (dense + sparse + ColBERT) → Qdrant (RRF fusion) → ColBERT reranking; grounded, sourced & cited generation

LLM: Mistral (sovereign, temp 0), OpenAI / Anthropic; prompt engineering, structured output, anti-hallucination & anti-injection guardrails, evaluation

Agents & orchestration: multi-agent workflows, MCP servers, Claude Code (Anthropic-certified) / Cursor tooling

Applied ML: on-device audio classification (YAMNet / TFLite, AudioSet), speech-to-text (Mistral)

AI infra: FastAPI, Qdrant, Docker, model self-hosting, Cloudflare Tunnel

PRODUCT STACK

Languages: Python, TypeScript, JavaScript, SQL, Java

Web: Next.js (App Router, RSC), React, Node.js, tRPC, PostgreSQL (Drizzle, Prisma), Stripe, Tailwind / shadcn

Reverse eng: Ghidra, Frida (proprietary protocols, encryption)

Ops: Docker / Compose, GitHub Actions (self-hosted runners), Linux (Debian), KVM, Nginx / Traefik

CERTIFICATIONS

Claude Code 101 ➤ · Claude Platform 101 ➤ · Claude 101 ➤ · Agent Skills
Anthropic Academy · 2026 · verifiable

EDUCATION

WebForce3 — Salesforce Developer (2022)

Simplon.co — Web developer (2017–2018)

Le 101 · IPI — programming / IT (2016–2017)

LANGUAGES

French: native · English: professional (technical)

SUMMARY

Senior full-stack engineer, ex-Accenture, focused on **applied AI**. I design and **run my own AI products in production**: a sovereign RAG assistant that can't make things up, LLM pipelines (STT + structured journaling), and on-device ML. I **ship fast with AI tooling** — Claude Code, Cursor, multi-agent orchestration — but I deliver **production-grade**: I understand the code underneath, I test, I secure, I deploy. From idea to a product that runs itself.

AI PRODUCTS IN PRODUCTION

JLLM — sovereign RAG assistant github.com/FelixDaCraft/JLLM ➤
Sourced & cited answers, ~99% faithfulness, 7/7 on injection tests · BGE-M3 → Qdrant (RRF) → reranking → Mistral

Ringdo — LLM call journaling github.com/FelixDaCraft/RingDo ➤
Native Android → backend · Mistral STT + LLM structuring (summary, topics, parties) + full-text search

BabyMonitor — on-device ML + reverse eng github.com/FelixDaCraft/v308-rtsp-hack ➤
ML cry detection (YAMNet TFLite) over an IP camera stream whose protocol I reverse-engineered (AES-128-ECB)

ALSO IN PRODUCTION

CupMetrics · Platinum CBD Cup [platinum-cbd-cup](https://platinum-cbd-cup.com) ➤
Multi-tenant event SaaS · Stripe · jury anonymization · 3D Three.js

Psychologist practice site [psy-cabinet](https://psy-cabinet.com) ➤
Client site delivered end-to-end · local SEO · self-service content back-office

EXPERIENCE

Freelance — Applied AI & Full-stack Aug 2023 → now
Design & operation of 5 products in production · Next.js / Node / Python / PostgreSQL

Accenture France — Nantes May 2022 – Jul 2023
Salesforce Developer · permanent · large-account project (Apex, LWC, Flows)

SQLI · Francecom · Yunow · 420 Green Road 2019 – 2022
Java/JEE (La Banque Postale), full-stack, Symfony, e-commerce (dropshipping Dolibarr/WordPress)

AI PROJECTS – DETAIL

PROJECT	JLLM – sovereign RAG assistant – github.com/FelixDaCraft/JLLM ↗
Problem	An LLM that can't lie about a corpus: every answer must be generated only from the source text , sourced and verifiable.
What I built	➤ Hybrid retrieval: BGE-M3 embeddings (dense + sparse + ColBERT) → Qdrant, RRF fusion, then ColBERT reranking ➤ Generation grounded by Mistral (sovereign, temperature 0), cited answers (chapter · section) ➤ Guardrails: systematic refusal of out-of-corpus queries, anti-injection (7/7 on tests), ~99% measured faithfulness ➤ FastAPI + chat UI, containerized with Docker, exposed via Cloudflare Tunnel, self-hosted
Stack	Python · FastAPI · Qdrant · BGE-M3 · Mistral · Docker · Cloudflare Tunnel
PROJECT	Ringdo – LLM call pipeline – github.com/FelixDaCraft/RingDo ↗
Problem	Automatically journal every phone call, from raw audio to an exploitable structured record.
What I built	➤ Native Android app (Java) that watches the call-recording folder and uploads each file ➤ Mistral speech-to-text on the audio, then LLM structuring (date, duration, parties, summary, topics) as structured output ➤ Full-text search over history · Node + PostgreSQL backend, Docker
Stack	Native Android · Java · Node.js · PostgreSQL · Mistral (STT + LLM) · Docker
PROJECT	BabyMonitor – on-device ML + reverse engineering – github.com/FelixDaCraft/v308-rtsp-hack ↗
Problem	Monitor a baby in real time from an IP camera locked to its proprietary cloud, with a smart alert on crying.
What I built	➤ Reverse engineering of the protocol (Ghidra / Frida): AES-128-ECB encryption, H.264 / HEVC decoding ➤ On-device ML cry detection: YAMNet TFLite (AudioSet audio classification) ➤ Self-hosted PWA + Web Push that pierces iOS silent mode · Node, Docker, Traefik
Stack	Python · YAMNet TFLite · Node Express · go2rtc / ffmpeg · PWA Web Push · Docker
METHOD	AI-assisted engineering (production-grade "vibe coding")
Day to day	➤ Development driven by Claude Code / Cursor: fast prototyping, then hardening (strict types, Vitest/Playwright tests, review) ➤ Multi-agent orchestration: specialized agent workflows (design, dev, review) and home-grown MCP servers to plug LLMs into my tools/data ➤ Delivery chain: one push = typecheck, tests, Docker build, migration, deploy, healthcheck (GitHub Actions CI/CD, self-hosted runners) ➤ Belief: AI makes you fast, but the value is shipping maintainable and secure – not accepting the output blindly
Availability	Remote · anywhere in France · employee or freelance depending on the role