Yann Maillot

APPLIED AI & FULL-STACK ENGINEER · RAG — LLM — AGENTS — ON-DEVICE ML

CONTACT

☎ +33 6 26 73 55 97

✉ yann.maillot.29@gmail.com

⌖ Remote · anywhere in France

in /in/yann-maillot

gh @FelixDaCraft

↗ yann.aynn.fr

AI SKILLS

RAG: hybrid pipeline BGE-M3 (dense + sparse + ColBERT) → Qdrant (RRF fusion) → ColBERT reranking; grounded, sourced & cited generation

LLM: Mistral (sovereign, temp 0), OpenAI / Anthropic; prompt engineering, structured output, anti-hallucination & anti-injection guardrails, evaluation

Agents & orchestration: multi-agent workflows, MCP servers, Claude Code (Anthropic-certified) / Cursor tooling

Applied ML: on-device audio classification (YAMNet / TFLite, AudioSet), speech-to-text (Mistral)

AI infra: FastAPI, Qdrant, Docker, model self-hosting, Cloudflare Tunnel

PRODUCT STACK

Languages: Python, TypeScript, JavaScript, SQL, Java

Web: Next.js (App Router, RSC), React, Node.js, tRPC, PostgreSQL (Drizzle, Prisma), Stripe, Tailwind / shadcn

Reverse eng: Ghidra, Frida (proprietary protocols, encryption)

Ops: Docker / Compose, GitHub Actions (self-hosted runners), Linux (Debian), KVM, Nginx / Traefik

CERTIFICATIONS

Claude Code 101 ↗ · Claude 101 ↗ · Agent Skills

Anthropic Academy · 2026 · verifiable

EDUCATION

WebForce3 — Salesforce Developer (2022)

Simplon.co — Web developer (2017–2018)

Le 101 · IPI — programming / IT (2016–2017)

LANGUAGES

French: native · English: professional (technical)

SUMMARY

Senior full-stack engineer, ex-Accenture, focused on **applied AI**. I design and **run my own AI products in production**: a sovereign RAG assistant that can't make things up, LLM pipelines (STT + structured journaling), and on-device ML. I **ship fast with AI tooling** — Claude Code, Cursor, multi-agent orchestration — but I deliver **production-grade**: I understand the code underneath, I test, I secure, I deploy. From idea to a product that runs itself.

AI PRODUCTS IN PRODUCTION

JLLM — sovereign RAG assistant github.com/FelixDaCraft/JLLM ↗

Sourced & cited answers, ~99% faithfulness, 7/7 on injection tests · BGE-M3 → Qdrant (RRF) → reranking → Mistral

Ringdo — LLM call journaling github.com/FelixDaCraft/RingDo ↗

Native Android → backend · Mistral STT + LLM structuring (summary, topics, parties) + full-text search

BabyMonitor — on-device ML + reverse eng github.com/FelixDaCraft/v308-rtsp-hack ↗

ML cry detection (YAMNet TFLite) over an IP camera stream whose protocol I reverse-engineered (AES-128-ECB)

ALSO IN PRODUCTION

CupMetrics · Platinum CBD Cup platinum-cbd-cup ↗

Multi-tenant event SaaS · Stripe · jury anonymization · 3D Three.js

Psychologist practice site psy-cabinet ↗

Client site delivered end-to-end · local SEO · self-service content back-office

EXPERIENCE

Freelance — Applied AI & Full-stack Aug 2023 → now

Design & operation of 5 products in production · Next.js / Node / Python / PostgreSQL

Accenture France — Nantes May 2022 – Jul 2023

Salesforce Developer · permanent · large-account project (Apex, LWC, Flows)

SQLI · Francecom · Yunow · 420 Green Road 2019 – 2022

Java/JEE (La Banque Postale), full-stack, Symfony, e-commerce (dropshipping Dolibarr/WordPress)

AI PROJECTS — DETAIL

PROJECT — **JLLM — sovereign RAG assistant** — github.com/FelixDaCraft/JLLM ↗

Problem — An LLM that can't lie about a corpus: every answer must be **generated only from the source text**, sourced and verifiable.

What I built

- Hybrid retrieval: BGE-M3 embeddings (dense + sparse + ColBERT) → Qdrant, **RRF** fusion, then ColBERT reranking
- Generation grounded by **Mistral** (sovereign, temperature 0), cited answers (chapter · section)
- Guardrails: systematic refusal of out-of-corpus queries, **anti-injection (7/7 on tests)**, ~99% measured faithfulness
- **FastAPI** + chat UI, containerized with Docker, exposed via Cloudflare Tunnel, self-hosted

Stack — Python · FastAPI · Qdrant · BGE-M3 · Mistral · Docker · Cloudflare Tunnel

PROJECT — **Ringdo — LLM call pipeline** — github.com/FelixDaCraft/RingDo ↗

Problem — Automatically journal every phone call, from raw audio to an exploitable structured record.

What I built

- **Native Android** app (Java) that watches the call-recording folder and uploads each file
- **Mistral speech-to-text** on the audio, then **LLM structuring** (date, duration, parties, summary, topics) as structured output
- Full-text search over history · Node + PostgreSQL backend, Docker

Stack — Native Android · Java · Node.js · PostgreSQL · Mistral (STT + LLM) · Docker

PROJECT — **BabyMonitor — on-device ML + reverse engineering** — github.com/FelixDaCraft/v308-rtsp-hack ↗

Problem — Monitor a baby in real time from an IP camera locked to its proprietary cloud, with a smart alert on crying.

What I built

- **Reverse engineering** of the protocol (Ghidra / Frida): AES-128-ECB encryption, H.264 / HEVC decoding
- On-device ML cry detection: **YAMNet TFLite** (AudioSet audio classification)
- Self-hosted PWA + Web Push that pierces iOS silent mode · Node, Docker, Traefik

Stack — Python · YAMNet TFLite · Node Express · go2rtc / ffmpeg · PWA Web Push · Docker

METHOD — **AI-assisted engineering (production-grade "vibe coding")**

Day to day

- Development driven by **Claude Code / Cursor**: fast prototyping, then hardening (strict types, Vitest/Playwright tests, review)
- **Multi-agent orchestration**: specialized agent workflows (design, dev, review) and home-grown **MCP servers** to plug LLMs into my tools/data
- Delivery chain: **one push = typecheck, tests, Docker build, migration, deploy, healthcheck** (GitHub Actions CI/CD, self-hosted runners)
- Belief: AI makes you fast, but the value is **shipping maintainable and secure** — not accepting the output blindly

Availability — Remote · anywhere in France · employee or freelance depending on the role